

Forecasting the black Sigatoka development rate: A comparison of machine learning techniques

Luis-Alexander Calvo-Valverde^{a,1}, Mauricio Guzmán-Quesada^b,
José-Antonio Guzmán-Alvarez^b, Pablo Alvarado-Moya^c

^a*DOCINADE, Instituto Tecnológico de Costa Rica, Computer Research Center,
Multidisciplinar program eScience, Cartago, Costa Rica*

^b*Dirección de Investigaciones, Corporación Bananera Nacional S. A., Guápiles, Costa Rica*

^c*DOCINADE, Instituto Tecnológico de Costa Rica, Cartago, Costa Rica*

Abstract

Pending.

Keywords: Machine learning, Black Sigatoka, Support vector regression,
Banana disease prediction, Biological warning system

1. Introduction

The black Sigatoka disease caused by the fungus *Mycosphaerella fijiensis* Morelet is the major pathological problem of banana and plantain crops in Central America, Panama, Colombia and Ecuador, as well as in many parts of Africa and Asia (Marín Vargas and Romero Calderón, 1995).

This disease attacks the plant leaves producing a rapid deterioration of the leaf area. It affects the growth and productivity of the plants due to the impairment of the photosynthetic process. Furthermore, it causes a reduction in the quality of the fruit, and promotes premature ripening of bunches, which is the major cause of product losses associated with the black Sigatoka.

For these reasons, warning systems have been developed to detect the disease and monitor its progress. For instance, the early warning system developed by Ganry and Meyer (1983) and modified by Ganry and Laville (1972) for the

Email address: lcalvo@itcr.ac.cr (Luis-Alexander Calvo-Valverde)

¹Corresponding author. (506)70104420

14 control of the yellow Sigatoka in Cameroon, was later adapted by Ternesien
15 (1985) and Fouré (1988) for the black Sigatoka.

16 This biological warning system is based on weekly observations of the disease
17 progression on young leaves of the plant. Figure 1 shows an example of three
18 progressive stages of the black Sigatoka.




Figure 1: Examples of three disease stages of the black Sigatoka. (a) Initial stage. (b) Intermediate stage, and (c) Advanced stage.

19 The disease progression is then quantified according to Fouré’s scale of the
20 symptom stages (Fouré, 1988) by means of numeric coefficients that describe
21 the degree of incidence and the severity of the disease development. These
22 coefficients are then used to calculate two variables: gross sum and state of
23 evolution.

24 The gross sum is based on the present disease progression stage and the
25 numeric coefficients, which increase with the progression of the symptoms and
26 the juvenility of the leaf. The state of evolution is calculated using the gross sum
27 and the foliar emission period. Decades ago threshold levels on these variables
28 were used as a guide to plan the spray schedules. Nowadays the fluctuation of
29 these two variables seems to better suggest appropriate times to spray (Marín
30 et al., 2003).

31 In Costa Rica the black Sigatoka is frequently treated with chemical fungi-
32 cides. Depending on the zone of production and the weather conditions, 45–
33 55 cycles/year of fungicide applications are required to keep this disease under
34 control and to produce the expected fruit quality for the international mar-

35 kets. This represents a cost per hectare per year in the range from US\$1600
36 to US\$2000; about 0.64–0.80 cents of the production costs for a 18.14 kg box,
37 which overall corresponds to 10%–12% of the total production costs.

38 The past and present rates of disease development can in principle be used
39 to predict its future behavior and to determine if a particular fungicide spray
40 program will be able to effectively treat the disease in an economically affordable
41 way (Chuang and Jeger, 1987). Phytopathological studies point out that climate
42 has a major effect on the development of the black Sigatoka, where the main
43 variables affecting it are precipitation, temperature, relative humidity and wind
44 (Marín Vargas and Romero Calderón, 1995). It can be expected that patterns in
45 these variables correlate with the disease development and hence its automated
46 discovery can support decision-making in the control of crop diseases.

47 In this work, we compare five machine learning techniques to predict the de-
48 velopment rate of the black Sigatoka disease: support vector regression (SVR),
49 echo state networks (ESN), ridge regression, elastic-net regression and ordinary
50 least squares linear regression, using input variables such as: air temperature,
51 relative humidity, solar radiation, precipitation and wind speed LISTO - completar las variables de entrada usa
52 to predict the evolution stage in the biological warning system LISTO - completar las variables de salida - Modi

53 The main contribution of this work² is the selection of the best machine
54 learning technique to forecast the black Sigatoka development rate according to
55 the metric indicated below.

56 The outline of the paper is as follows: Section 2 presents related works and
57 Section 3 summarizes the machine learning techniques selected for the analy-
58 sis. In Section 4 we present the methodology used in this study and describe
59 data used for its verification. The results and their discussion are presented
60 in section 5. The Section 6 concludes this article and presents lines for future

²Definitivamente ese no es el aporte principal del artículo. Cuando esté listo el resto hay que volver aquí. Creo que va a ser algo como mostrar la relevancia de la etapa de preprocesamiento, o mostrar que a pesar de ser sistemas complejos, con variables caóticas, lo mejor en este caso son predictores lineales, o algo por el estilo - Pendiente

61 works.

62 2. Related works

63 Several efforts have been made to apply machine learning techniques in the
64 automated discovery of relationships between environmental variables and quan-
65 tified descriptors for variables of agricultural interest such as the progress of
66 diseases. Huang et al. (2010) summarize in their survey the development of soft
67 computing techniques in agricultural and biological engineering, especially in the
68 soil and water context for crop management and decision support in precision
69 agriculture, including fuzzy logic, artificial neural networks, genetic algorithms,
70 Bayesian inference and decision trees. They do not present numeric results of
71 each paper, only mention the main idea. Similarly, Kim et al. (2014) survey
72 more recent prediction methods for crop pests using regression and machine
73 learning approaches. Nor do they provide numerical results of each paper.

Please review:

74 MODIFICADO - Falta indicar qué concluyen los dos surveys anteriores — como indico, en
realidad ellos mencionan qué métodos han usado en los paper, pero no dan R^2 o rmse, ni
detallan resultados en general

75 In general, the machine learning methods applied to predict the evolution
76 of plant diseases, can be classified in two main approaches: 1) those whose
77 main inputs are images, and 2) Those whose main inputs are environmental
78 and biological variables. Our study focuses in the second case.

79 Romero Calderón (1995) relied on regression models using a stepwise proce-
80 dure to predict incubation and disease latency periods for the black Sigatoka.
81 He collected environmental data from two different farms in Costa Rica between
82 December 1993 and August 1995. The prediction models reached coefficients of
83 determination R^2 of 69% or 78% on the observed data for the incubation and
84 disease latency periods, respectively; however, the cross validation on indepen-
85 dent data sets failed. In contrast, our proposal presents a model that can be
86 generalizable to other farms that have data.

87 More recently, Glezakos et al. (2010) used genetic algorithms (GA) and neu-
88 ral networks (NN) to identify the Tobacco Rattle Virus (TRV) and the Cu-
89 cumber Green Mottle Mosaic Virus (CGMMV). The method was tested against
90 some of the most commonly used classifiers in machine learning (Bayes clas-
91 sifiers, decision trees and k -nearest neighbors) via cross-validation and proved
92 their applicability in these kind of problems. These authors do not prove their
93 methods in Sigatoka disease and they do classification. Instead we do regression.

94 Alves et al. (2011) used geoinformation techniques to develop predictive
95 models in the study of risk areas to soybean rust, coffee leaf rust, and banana
96 black Sigatoka, under consideration of Brazil's climatic characteristics and the
97 distribution of soybean, coffee and banana crops. Temperature and rainfall
98 data were acquired for the period from 1950 to 2000, and simulated data were
99 generated for 2020, 2050 and 2080 using the SRES A2 climate change scenarios.
100 Using principal components analysis, a single variable was generated as a linear
101 combination of 57 input variables, in order to determine an index explaining
102 87%, 88% and 90% of the data variability³ of soybean, coffee and banana crops,
103 respectively, in municipal districts across Brazil. The climatic model was used
104 to generate the zoning of the three plant diseases, using temperature and leaf
105 wetness as input. This methodology enabled the visualization of the changes in
106 areas favorable for epidemics under possible future scenarios of climate change.
107 How intermediate result, they characterized the monocyclic process of the black
108 Sigatoka using nonlinear regression. Although they do not present the detailed
109 results, it no seem that they wanted to predict the progression of the black
110 Sigatoka in one, two or more periods ahead, how we do.

111 Other applications of machine learning methods in precision agriculture in-
112 clude the use of support vector regression to predict carcass weight in beef cattle
113 in advance to the slaughter (Alonso et al., 2013), machine learning assessments
114 of soil drying for agricultural planning (Coopersmith et al., 2014), and early de-
115 tection and classification of plant diseases with support vector machines based

³variabilidad de qué? - MODIFICADO: de los datos de cada uno

116 on hyperspectral reflectance (Rumpf et al., 2010).

117 Furthermore, there have been attempts to generate software tools. Camargo
118 et al. (2012) presented an information system for the assessment of plant disor-
119 ders (Isacrodi). They showed that human experts will attain a much accurate
120 assessment than the Isacrodi classifier, particularly when provided with sam-
121 ples from the affected crop. However, in those cases where such expertise is not
122 available, the authors suggest that Isacrodi can still provide valuable support to
123 farmers. Isacordi includes 15 crop disorders, but the black Sigatoka is none of
124 them. The prediction process is based on multi-class support vector machines.

125 Regarding the prediction of the black Sigatoka disease development with
126 machine learning methods, Bendini et al. (2013) presented a study on the risk
127 analysis of its occurrence based on polynomial models. A case study was de-
128 veloped in a commercial banana plantation located in Jacupiranga, Brazil. It
129 was monitored weekly from February to December 2005. The data included the
130 weekly monitoring of the disease’s evolution stage, time series of meteorological
131 data and remote sensing data. They obtained a model to estimate the evolution
132 of the disease from satellite imagery. This model relates gray levels (NC) of
133 the band 2 images of the Landsat-5 satellite, with the progress status or disease
134 severity (EE). The authors claim to reach an R^2 of 90%.

135 There are also works related to the banana fruit. Soares et al. (2014) apply
136 two techniques: artificial neural networks (ANNs) and multiple linear regression
137 (MLR) in banana plant to predict the yield. Their results show that the neural
138 network is more accurate in forecasting the weight of the bunch in comparison
139 to the multiple linear regressors in terms of the mean prediction-error ($MPE =$
140 1.40), mean square deviation ($MSD = 2.29$) and coefficient of determination
141 ($R^2 = 91\%$).

142 Although these studies have their contribution, none proposed the kind of
143 preprocessing that we present, nor pose how to predict more than one period
144 ahead without trying to predict climate.

Please review:

MODIFICADO - Importantísimo es explicar aquí brevemente por qué esos otros trabajos no se aplican al caso de la Sigatoka, o qué es lo que han hecho mal, que este trabajo sí hace bien

3. Compared regression techniques

In the prediction of the development rate of the black Sigatoka, we compare techniques such as least squares or ridge regression, commonly encountered in the agricultural literature with machine learning methods such as support vector regression, elastic regression and echo state networks, where the parameter space of each technique is also taken into account.

3.1. Ordinary least squares regression

Given a data set

$$D = \{(\mathbf{x}_i, y_i) \mid i = 1 \dots n\} \quad (1)$$

composed of the d -dimensional⁴ feature vectors $\mathbf{x}_i \in \mathbb{R}^d$ and the corresponding responses y_i . The ordinary least squares regression (OLSR) fits a linear model $\tilde{y}_i = f(\mathbf{x}_i) = \langle \mathbf{w}, \mathbf{x}_i \rangle$ such that the sum of squares of the residuals $(\tilde{y}_i - y_i)$ is minimized. Let \mathbf{X} be the $n \times d$ feature matrix containing the i -th data sample \mathbf{x}_i^T in its i -th row and \mathbf{y} contain all the responses y_i corresponding to each row, then the least squares regression finds

$$\hat{\mathbf{w}} = \arg \min_{\mathbf{w}} E(\mathbf{w}) \quad (2)$$

with the error function

$$E(\mathbf{w}) = \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2$$

The solution is found by means of the pseudoinverse $\hat{\mathbf{w}} = \mathbf{X}^+ \mathbf{y} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$ or equivalently by the singular value decomposition of \mathbf{X} (Press et al., 2007).

⁴Without loss of generality assume that the first component of every vector \mathbf{x}_i is always 1.

Please review:

Alex, no sé hasta donde hiciste en las pruebas ese “arreglo” de que los datos tengan todos una entrada igual a 1. De no ser así, el método no puede encontrar el offset, y sería como forzar que el modelo tenga que pasar por cero... a menos por supuesto que la constante se haya sacado de la ecuación al optimizar...

COMENTARIO: Modificado según lo conversado

3.2. Ridge regression

In contrast to the OLSR, for the ridge regression (RR) Hoerl and Kennard (1988) proposed to add a term to penalize large weights into the error function

$$E(\mathbf{w}) = \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 + \alpha \|\mathbf{w}\|_2^2$$

where the parameter $\alpha > 0$ controls how strong is the shrinking of the estimates towards zero. This shrinkage introduces some bias but helps to reduce the variance of the estimate. The solution of the optimization problem (2) in this case is given by $\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$.

3.3. Elastic net regression

Instead of L_2 regularization prior ($\alpha \|\mathbf{w}\|_2^2$) included in the ridge regression, Tibshirani (1996) used an L_1 term ($\lambda \|\mathbf{w}\|_1$) for his lasso estimator, which permits to select a subset of the available features by zeroing the weights of the deselected features. If the dimension d of the data is larger than the number n of data samples, lasso will select a maximum of d variables.

The elastic net regression (ENR) of Zou and Hastie (2005) combines both L_1 and L_2 priors of the ridge and lasso estimators such that the error function is now

$$E(\mathbf{w}) = \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 + \alpha \|\mathbf{w}\|_2^2 + \lambda \|\mathbf{w}\|_1$$

This combination of priors still allows to learn a sparse model with only a few weights being non-zero like in the case of lasso, but still maintaining the regularization properties of the ridge regression (Pedregosa et al., 2011).

The elastic net is useful when multiple features are correlated: lasso will likely pick one of these at random, while the elastic net will still likely pick both.

Please review:

Algo que no entiendo es que las elastic net son en realidad una generalización del ridge, Lasso y OLSR. Así que en los experimentos con las selección adecuada de parámetros este método debería comportarse al menos igual o mejor que esos otros regresores! No sé hasta donde sea justificable usar este Y los otros, porque este ES los otros...

COMENTARIO: Revisando la documentación de la biblioteca (ver abajo), parece que por implementación no debería excluirse OLSR por separado. Estos son dos parámetros de las ElasticNet `alpha` : float. Constant that multiplies the penalty terms. Defaults to 1.0 See the notes for the exact mathematical meaning of this parameter. `alpha = 0` is equivalent to an ordinary least square, solved by the LinearRegression object. For numerical reasons, using `alpha = 0` with the Lasso object is not advised and you should prefer the LinearRegression object. `l1ratio` The ElasticNet mixing parameter, with `0 ≤ l1ratio ≤ 1`. For `l1ratio = 0` the penalty is an L2 penalty. For `l1ratio = 1` it is an L1 penalty. For `0 ≤ l1ratio ≤ 1`, the penalty is a combination of L1 and L2

Please review:

Alex, veo que usas mucho (Pedregosa et al., 2011) para los métodos, pero ese es solo el paper de un tool y no los proponentes originales de los métodos. Usualmente uno hace referencia a algún artículo, libro o tutorial donde ojalá los que propusieron el método son los que lo explican. Ahí metí entonces en las referencias otros artículos por ese motivo.

COMENTARIO: muchas gracias..

3.4. Support Vector Regression (SVR)

Please review:

Estoy seguro que los lectores van a solicitar reducir esto y dejar solo la referencia. Dejé solo lo que considero relevante, pero podés reducirlo más, si querés.

COMENTARIO: Ok, gracias

From the perspective of Support Vector Regression (SVR) the regression function is usually formulated as

$$\tilde{y} = f(\mathbf{x}) = \langle \mathbf{w}, \mathbf{x} \rangle + b \quad (3)$$

The weights are selected in a convex optimization problem (Smola and Schölkopf,

2004):

$$\begin{aligned} & \text{minimize} \quad \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n (\xi_i + \xi_i^*) \\ & \text{subject to} \quad \begin{cases} y_i - \langle \mathbf{w}, \mathbf{x}_i \rangle - b & \leq \epsilon + \xi_i \\ \langle \mathbf{w}, \mathbf{x}_i \rangle + b - y_i & \leq \epsilon + \xi_i^* \\ \xi_i, \xi_i^* & \geq 0 \end{cases} \end{aligned}$$

178 where ϵ is the maximal allowed deviation of the targets \tilde{y}_i from the responses y_i ,
 179 the slack variables ξ_i and ξ_i^* allow to cope with otherwise unfeasible constraints
 180 for the optimization problem, and the constant $C > 0$ controls the trade-off
 181 between the flatness of f and the tolerance to deviations larger than ϵ .

182 Note that since OLSR, RR and ENR use a squared error function, data
 183 outliers will have a strong influence on the resulting weights \mathbf{w} . On the SVR
 184 formulation, however, the usage of the L_1 norm and the slack variables consid-
 185 erably restrict or completely block the influence of those outliers.

The SVR problem is reformulated by means of the dual optimization problem into (Smola and Schölkopf, 2004)

$$\mathbf{w} = \sum_{i=1}^n (\alpha_i - \alpha_i^*) \mathbf{x}_i \quad (4)$$

where $\alpha_i, \alpha_i^* \in [0, C]$ are Lagrange multipliers subject to $\sum_{i=1}^n (\alpha_i - \alpha_i^*) = 0$. In this so-called *Support Vector expansion* the weights are expressed as a linear combination of the data set patterns \mathbf{x}_i . Inserting (4) in (3) leads to

$$f(\mathbf{x}) = \sum_{i=1}^n (\alpha_i - \alpha_i^*) \langle \mathbf{x}_i, \mathbf{x} \rangle + b \quad (5)$$

186 The Lagrange multipliers α_i, α_i^* are both non-zero only for those data points
 187 where $|f(\mathbf{x}_i) - y_i| \geq \epsilon$. Hence, the expansion of \mathbf{w} in terms of \mathbf{x}_i is sparse.
 188 Those data points with non-vanishing coefficients are called *Support Vectors*
 189 (Wei et al., 2013).

190 Additionally, in (5) it is possible to employ the *kernel trick* and replace the
 191 terms $\langle \mathbf{x}_i, \mathbf{x} \rangle$ with the evaluation of any Mercer kernel $k(\mathbf{x}_i, \mathbf{x}) = \langle \phi(\mathbf{x}_i), \phi(\mathbf{x}) \rangle$,

where $\phi(\mathbf{x})$ is a non-linear mapping of the input space onto a higher (even infinite) dimensional feature space. The kernel evaluation draws unnecessary the explicit evaluation of the non-linear mapping, and it allows to solve non-linear regressions in the input space by implicitly mapping the samples through the kernel into the higher dimensional space, where the linear regression occurs (Alonso et al., 2013).

Kernels used in this work were:

Linear kernel:

$$K(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i^T \mathbf{x}_j$$

Radial basis function (RBF):

$$K(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(-\frac{\|\mathbf{x}_i - \mathbf{x}_j\|^2}{2\sigma^2}\right)$$

where σ is the parameter of gaussian model.

Sigmoid:

$$K(\mathbf{x}_i, \mathbf{x}_j) = \tanh\left[-c + \frac{\mathbf{x}_i \mathbf{x}_j}{\sigma^2}\right]$$

with $c \geq 0$ and σ^2 is the scaling vector.

Please review:

LISTO - Habría que insertar aquí los kernels utilizados en los experimentos

3.5. Echo State Networks (ESN)

Recurrent neural networks (RNN) are capable of learning temporal patterns by feeding neuron outputs back into lower layers. Their training (usually by means of error backpropagation) is in general slow. Echo state networks (ESN) are a particular type of recurrent neural network with a sparsely connected random hidden layer where only the weights of the output neurons are changed at training. The randomly selected weights at the input and middle layers (called *reservoir*) reproduce temporal patterns (*echoes*) that the output layer learns to select during the training (Lukoševičius and Jaeger, 2009).

For a given training input signal $u(n) \in \mathbb{R}^{N_u}$ a desired target output signal $y^{target}(n) \in \mathbb{R}^{N_y}$ is known. Here $n = 1, \dots, T$ is the discrete time and

213 T is the number of data points in the training dataset.⁵

214 The training seeks to learn a model with output $y(n) \in \mathbb{R}^{N_y}$, where $y(n)$
215 matches $y^{target}(n)$ as close as possible, by means of the minimization of an
216 error measure $E(y, y^{target})$ such that it also generalizes well to unseen data
217 (Lukoševičius, 2012).

218 4. Specification of data and methodology

219 Since the suitability of a machine learning technique to a particular problem
220 is entirely depend on the nature of the data, we describe in this section, first,
221 the data set employed in the study, followed by the methodology to compare
222 the chosen techniques under consideration of their parameter space.

223 4.1. Data

224 The data used for the current study was acquired in two research farms
225 of Corbana in Costa Rica⁶ : *28 Millas* located in the region of Matina, and
226 *La Rita* located in Pococí, both in the province of Limón, Costa Rica. Both
227 farms produce banana fruit *Musa* sp. AAA group ‘Grande Naine’ (Cavendish
228 subgroup).

229 The available input and output variables are summarized in Table 1.

230 The data was captured for La Rita between the 48th week of 2002 to the
231 17th week of 2015 (647 weeks); for 28 Miles the data was captured between the
232 37th week of 2003 and the 18th week of 2015 (605 weeks). The data on the
233 biological warning system were collected once a week.

234 The meteorological stations of Corbana acquire data every five minutes.

235 However, for the current study weekly data were used.⁷ computed on the data
236 collected by nearby stations in each farm. Experiments were carry out with daily

⁵No puede ser! – Así está en el paper original, lo puse en el repositorio

⁶Both farms were also use in the study of Romero Calderón (1995). Back then, *La Rita* was referred to as *Waldeck*.

⁷Hay que explicar por qué? → reducción de ruido, etc.

Symbol	Description	Units
$T_{a_{max}}$	Maximal air temperature	[°C]
$T_{a_{min}}$	Minimal air temperature	[°C]
\bar{T}_a	Mean air temperature	[°C]
\bar{H}	Mean relative humidity	[0 – 100]
H_{min}	Minimal relative humidity	[0 – 100]
H_{max}	Maximal relative humidity	[0 – 100]
\bar{R}	Mean solar radiation	[W/m ²]
P	Precipitation	[mm]
W_{max}	Maximal wind speed	[m/s]
\bar{W}	Mean speed wind	[m/s]
E_s	Biological warning system – Evolution Stage	> 0

Table 1: Variables available for the learning algorithms

periodicity in meteorological variables and the results proved do not improve the prediction. Besides, weekly data pretend to diminish noise due sensor accuracy, missing values and outliers no detected.

Please review:

Alex, Falta explicar cada variable y sus unidades y rangos usuales, lo que es relevante luego cuando se hable de la normalización.

Humidity ¿cuál se usa? ¿la relativa? ¿la absoluta?, ¿la específica? Las otras unidades las supuse, así que hay que revisar que estén bien!

No tengo idea qué es el “Sum precipitation” y en qué unidades estaría. Parece no ser algo estándar o medido directamente, ¿o sí? así que amerita una explicación aquí.

Dijiste que las variables se promedian para una semana. ¿Qué significan entonces las variables que dicen “mean”?

COMENTARIO: Modifiqué estas explicaciones para que quedara más claro el detalle de las variables. Me dirás qué te parece, si las dejamos así o amplío.

Please review:

The mean variables ...

The value to be predicted in all cases is the evolution stage E_s , which is a measure of the level of disease progression.

244 4.2. Data preprocessing

245 Data taken on real farms during more than a decade is expected to contain
246 outliers, noise and missing samples. These problems are caused by human errors
247 or by technical defects on the instruments used. In the preprocessing step
248 described in this section these problems need to be detected and fixed before
249 moving them to the next processing stages.

250 In the farm 28 Miles 1% and in La Rita 2.25% of the data were missing.
251 To fill-in the missing values spline interpolation⁸ was used. The data collected
252 did not exhibit outliers.

253 Due the fact that the variables measure meteorological or biological pro-
254 cesses, they are discretized in order to reflect data trends⁹. The variation range
255 of each variable is uniformly discretized. This discretization removes the effect
256 of small variations in the data, either by inaccuracies of the instruments (me-
257 teorological variables) or by subjective bias introduced by the human collecting
258 the data (biological warning system).

259 The coefficient of variation $C_v(x) = \sqrt{E((x - E(x))^2)}/E(x)$ of each variable
260 x is used to determine the number n of discretization levels as $n = \lfloor 100 C_v(x) \rfloor$,¹⁰

⁸falta una referencia

⁹Formalmente la discretización NO captura tendencias, sino la derivada de la función en el tiempo. Para que no se capture ruido uno simplemente usa un filtro pasabajas primero (los estadísticos le dicen diferente, pero no sé cómo). En otras palabras, si querés la tendencia, con usar la derivada de un filtro gaussiano sobre los datos ya se ve la tendencia muy bien con el signo de esa derivada. Con la varianza del filtro uno ajusta qué nivel de detalle le interesa, o lo qué es lo mismo, cuanto tiempo quiere usar para el pronóstico de tendencias (variable crece, está igual o decrece). El valor de salida de la derivada filtrada también se puede discretizar si solo interesa esa tendencia de forma difusa, pero quizá para la regresión es mejor mantenerla continua (si está filtrada, con un filtro o mejor aún con un banco de filtros, al estilo wavelets).

Eso es TOTALMENTE independiente de la cuantificación. Ahora, un cuantificador tiene en principio un efecto de filtro pasabajas, pero eso reduce el ruido y en realidad no dice nada de tendencias. En resumen, formalmente la cuantificación en sí no refleja nada de un comportamiento temporal, y “tendencia” es en principio un comportamiento temporal!

¹⁰Hay algo que no entiendo. ¿De dónde sale el 100? ¿Por qué 100? Así como está esto, el mismo rango de variación de temperatura produce diferentes n , con solo un cambio en el

261 where $\lfloor \cdot \rfloor$ is the round operator.

262 Each feature **was scaled**¹¹ to fit in the range 0 and 1. The variable E_s to
263 be predicted was not scaled.

Please review:

264 Aquí faltan la fórmula general de normalización usada. Lo que sigue lo estoy asumiendo, pero
puede ser que no sea así

265 Each variable $x \in [x_{\min}, x_{\max}]$ was normalized into the interval $[0, 1]$ with
266 the linear map $x_n = mx + b$ with $m = 1/(x_{\max} - x_{\min})$ and $b = -mx_{\min}$.

Please review:

267 Explicar brevemente la razón de la normalización, ojalá con una referencia bibliográfica

268 4.3. Evaluation criteria

269 Although there are many types of indicators to assess the quality of the
270 prediction, here the coefficient of determination (R^2) and the Root Mean Square
271 Error ($RMSE$). This decision is supported by the widespread use of the former
272 indicator in the agriculture and the latter in machine learning (Soares et al.,
273 2013, 2014; Ibrahim and Wibowo, 2014; Demir and Bruzzone, 2014).

Given n records y_i , $i = 1 \dots n$ of the actual outcome of a process. The mean
 \bar{y} of the observed data is given by

$$\bar{y} = \frac{1}{n} \sum_{i=1}^n y_i$$

Let \hat{y}_i be the predicted value for y_i . Then, the mean square error (MSE) S_e^2

promedio! Eso es muy extraño. Por otro lado, los rangos de las variables son tan diferentes, que el número n va a terminar siendo casi cualquier cosa. Yo hubiese esperado que la desviación estándar se use para discretizar las variables. Otro problema formal más serio, es que el coeficiente de variación solo se puede usar/tiene sentido si la escala de medición es de razón, y la temperatura, la precipitación y la humedad son usualmente mediciones de intervalo! Los L_i ni tan siquiera sé qué son porque creo que son algo subjetivos a la persona que los toma ¿no?, aunque creo que iban con el porcentaje de área afectado, lo que también sería medición de intervalo!

¹¹ ¿escaladas cómo? ¿sx? ¿sx + b? ¿sf(x)?

and the S_e^2 are estimated as

$$S_e^2 = \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{n} \quad S_R^2 = \frac{\sum_{i=1}^n (\hat{y}_i - \bar{y})^2}{n}$$

The root mean square error is defined as $RMSE = \sqrt{S_e^2}$ and the coefficient of determination is

$$R^2 = \frac{S_R^2}{S_R^2 + S_e^2}$$

Please review:

Favor revisar. Vi una representación en términos de $R^2 = 1 - MSE/\sigma^2$ (con σ^2 la varianza), que no sé si está bien o no, pero lo mejor será poner la fuente de donde se tomó la anterior definición. Había en el S_R^2 un \bar{y}_i pero el promedio no depende de i por lo que no tenía sentido. Supuse que era el promedio, pero no sé. El cuadrado en las S^2 ¿está bien?, o indirectamente estás diciendo que $RMSE = S_e$, que sería una mejor notación matemática (dejando RMSE solo como abreviatura textual)!

4.4. Programming environment

We use the Python programming language with the Integrated Development Environment (IDE)¹² Spyder (Analytics, 2015), particularly with the libraries Pandas (McKinney, 2010) and Numpy (van der Walt et al., 2011).

The implementation for SVR, ridge and ordinary least squares regressions in scikit-learn (Pedregosa et al., 2011) were used. Adjustments to the ESN implementation code of Lukoševičius (2012) were necessary to allow its integration into our experimental framework.

All experiments were performed on a PC computer with an Intel® Core i7-4800MQ processor, 2.70 GHz, 16.0 GB RAM, under the operating system MS Windows 8 Pro.

¹²El IDE es para la reproducción de resultados irrelevante, y por lo tanto información circunstancial que es mejor evitar. Lo relevante es lo que tenga efectos directos en los resultados. Si yo prefiero usar un editor de texto corriente para programar, el resultado final es el mismo, si uso el mismo compilador/intérprete. Eso es la información relevante: ¿qué versión de Python concretamente? ¿Se compiló el Python o se usó interpretado? ¿Cuáles versiones de Pandas y Numpy en concreto se usaron?

286 4.5. Methodology

287 The evaluation of the techniques under consideration of their parameter
288 space, was performed in two stages, described below.

289 4.5.1. Phase one

290 In the phase one, ten-fold-cross-validation on the total set of machine learn-
291 ing methods under a subset of configurations was evaluated:

- 292 • Patterns: $n \times m$, with $n = 1 \dots 8$ and $m = 1 \dots 2$.

293 **Please review:**

No se entiende que quiere decir eso!!

- 294 • Methods: support vector regression with the kernels linear, gaussian and
295 sigmoid; echo state networks; ordinary least squares linear regression, ridge
296 regression and elastic-net regression.

Please review:

297 Cada uno de esos métodos tiene a su vez un espacio paramétrico. Eso tiene que quedar
claro aquí: cómo se barrió el espacio paramétrico, para saber el nivel de detalle de la
prueba! Eso es importantísimo, porque si no no se puede decir nada respecto a los
métodos, sino solo respecto la configuración concreta probada!

- 298 • Variables included in the model:

- 299 – All variables.
300 – From the set $\{\overline{T}_a, \overline{H}, P, \overline{W}\}$ use the subsets with one, two or four
301 elements. These variables have the largest impact on the disease
302 development (Marín Vargas and Romero Calderón, 1995).

303 4.5.2. Phase two

304 In the second phase, the best configurations obtained in phase one are used
305 to validate with the last 50 and 100 weeks.

Please review:

306 Hmm, ¿entonces cuántas/cuáles semanas se usaron para la fase uno?
307

308 5. Results and discussion

309 In this section we present the main results for phase and their discussion.

310 Phase one

311 Figure 2 shows the best R^2 for each algorithm in the experiment. Results
 312 are group by farm. Though La Rita obtains different results in magnitude than
 313 28 Millas, the trend is similar. In both farms, the best results are for linear
 314 models, second position is occupied for Echo State Networks and SVR with
 315 gaussian and sigmoid kernels are the worst results. In linear models, to predict
 316 one week ahead is better than two weeks ahead, and this is better than three
 317 weeks ahead.




Figure 2: Phase one - Best R^2 for each algorithm

318 Figure 3 presents, for one, two and three weeks ahead, the best R^2 . Results

are group by farm. In general, to predict one week ahead is better than two weeks ahead and so on. The number of weeks consider in the observation range in the pattern is not the main discriminant factor, but it is clear that we get better R^2 for one week ahead than two weeks ahead and so on.




Figure 3: Phase one - Best R^2 for each observation range

Figure 4 shows the best R^2 for each variables combination. Results are group by farm. The better results are obtained with \bar{T}_a and the combination of \bar{T}_a with \bar{W} , in both farms of similarly. You can note that the use of all variables in the model or the inclusion of the four variables suggest for expert criteria do not improve significantly the results, then the use of more sensors do not assure a better result.




Figure 4: Phase one - Best R^2 for each variable combination

Figure 5 shows the Pareto frontier for each farm with respect to R^2 and $RMSE$. The Rita obtains upper R^2 with respect to 28 Millas, but 28 Millas obtains better $RMSE$ than La Rita. This situation arise because $RMSE$ considers errors only with respect the prediction and in 28 Millas the average of Stage of Evolution is 4316.16, unlike, in La Rita the average is 5507.30. So, in La Rita we obtains higher errors in absolute values. R^2 is a relative metric between 0 thru 1 and it is less sensitive to absolute values.




Figure 5: Phase one - Pareto frontier for R^2 and $RMSE$

336 The Pareto frontier for the La Rita farm is composed by 96 elements. The
 337 Table.2 shows the composition about variables and observation ranges.

Table 2: Composition of the Pareto frontier - La Rita - Phase one

Variable	Observation range	Quantity	Max R^2	Min $RMSE$
Pair \bar{T}_a \bar{W}	1 to 1	36	64.25%	714.51
	2 to 1	6	62.97%	695.10
All	1 to 1	18	62.98%	701.95
	2 to 1	12	61.76%	679.92
	3 to 1	6	60.60%	676.42
	5 to 1	2	60.37%	672.39
\bar{T}_a	1 to 1	12	63.60%	708.77
	2 to 1	4	62.23%	689.55

338 Similarly, the Pareto frontier for the 28 Millas farm is composed by 75 el-
 339 ements. The Table.3 shows the composition about variables and observation

340 ranges.

Table 3: Composition of the Pareto frontier - 28 Millas - Phase one

Variable	Observation range	Quantity	Max R^2	Min $RMSE$
Pair \overline{T}_a \overline{W}	1 to 1	8	57.80%	438.09
All	9 to 1	2	50.93%	397.93
	10 to 1	2	50.97%	398.81
	8 to 1	6	51.62%	398.93
	7 to 1	2	52.25%	400.28
	6 to 1	2	53.16%	404.14
	4 to 1	2	54.32%	407.54
\overline{T}_a	1 to 1	8	59.09%	439.44
Pair \overline{T}_a \overline{H}	1 to 1	8	57.51%	428.61
	2 to 1	20	56.91%	414.37
	3 to 1	3	54.41%	411.55
	4 to 1	3	53.34%	406.65
Pair \overline{T}_a P	3 to 1	9	56.23%	422.76

341 We can conclude that the best configuration in both farms is to consider
342 the climate and the evolution stage of the current week to predict the evolution
343 stage of the next week.

344 **Phase two**

345 In the second phase, the best configurations obtained in phase one were used
346 to validate with the last 50 and 100 weeks.

347 Figure 6 shows the best R^2 for each algorithm in the experiment. Results
348 are group by farm. Even if linear models continue with good R^2 , Echo State
349 Networks improve their scores because in 50 and 100 last weeks validation, we
350 are in presence of climate change, then the behaviour is less lineal.




Figure 6: Phase two - Best R^2 for each algorithm

Figure 7 presents, for one, two and three weeks ahead, the best R^2 . Results are group by farm. The results confirms that, in general, to predict one week ahead is better than two weeks ahead, two than three and so on, this for both farms.




Figure 7: Phase two - Best R^2 for each observation range

Figure 8 shows the best R^2 for each variables combination. Results are group
 by farm. This results confirm that \overline{T}_a and the combination of \overline{T}_a with \overline{W} , in
 both farms are the best variables combinations.




Figure 8: Phase two - Best R^2 for each variable combination

Figure 9 shows the Pareto frontier for each farm with respect to R^2 and $RMSE$. You can note that the behaviour of R^2 and $RMSE$ is similar to the phase one.




Figure 9: Phase two - Pareto frontier for R^2 and $RMSE$

361 In phase two, the Pareto frontier for the La Rita farm is composed by 2
 362 elements. The Table.4 shows the composition about variables and observation
 363 ranges.

Table 4: Composition of the Pareto frontier - La Rita - Phase two

Variable	Observation range	Quantity	Max R^2	Min $RMSE$
All	1 to 1	2	69.33%	353.33

364 In 28 Millas Farm, the Pareto frontier is composed by 1 element. The Table.5

365 details the result.

Table 5: Pareto frontier - 28 Millas - Phase two

Variable	Observation range	Quantity	Max R^2	Min $RMSE$
Pair \bar{T}_a P	1 to 1	1	61.36%	339.89

366 Again, similar to phase one, we can conclude that the best configuration in
 367 both farms is to consider the climate and the evolution stage of the current week
 368 to predict the evolution stage of the next week, one week to predict one week
 369 ahead with combinations of variables listed above.

370 6. Conclusions

371 6.1. Future work

372 Acknowledgements

373 The authors would like to thank Corbana S.A. for providing the data for
 374 this research.

375 References

- 376 Alonso, J., Rodríguez Castañón, Á., Bahamonde, A., 2013. Support vector re-
 377 gression to predict carcass weight in beef cattle in advance of the slaughter.
 378 Computers and Electronics in Agriculture, 116–120.
- 379 Alves, M., de Carvalho, L., Pozza, E., Sanches, L., Maia, J., 2011. Ecologi-
 380 cal zoning of soybean rust, coffee rust and banana black Sigatoka based on
 381 Brazilian climate changes. Procedia Environmental Sciences 6, 35–49.
- 382 Analitycs, C., 2015. Anaconda. Online; accessed 25-April-2016.
 383 URL <https://www.continuum.io/>
- 384 Bendini, H., Moraes, W., da S. Silva, Tezuka, E., Cruvinel, P., 2013. Análise de
 385 risco da ocorrência de sigatoka-negra baseada em modelos polinomiais: um
 386 estudo de caso. Tropical Plant Pathology 38 (1), 35–43.

387 Camargo, A., Molina, J., Cadena-Torres, J., Jiménez, N., Kim, J., 2012. Intelli-
388 gent systems for the assessment of crop disorders. Computers and Electronics
389 in Agriculture 85, 1–7.

390 Chuang, T., Jeger, M., 1987. Predicting the rate of development of black Sigatoka (*Mycosphaerella fijiensis* var. *difformis*) disease in southern Taiwan. Phytopathology 77, 1542–1547.

393 Coopersmith, E. J., Minsker, B. S., Wenzel, C. E., Gilmore, B. J., 2014. Machine
394 learning assessments of soil drying for agricultural planning. Computers and
395 Electronics in Agriculture 104, 93–104.

396 Demir, B., Bruzzone, L., 2014. A multiple criteria active learning method for
397 support vector regression. Pattern Recognition, 2558–2567.

398 Fouré, E., 1988. Stratégies de lutte contre la cercosporioses noire des bananiers
399 et plan-tains provoquée par *Mycosphaerella fijiensis* Morelet. L’avertissement
400 biologique au Cameroun. Evaluation des possibilités d’amélioration. Fruits
401 43 (5), 269–274.

402 Ganry, J., Laville, E., 1972. La lutte contrôlée contre le Cercospora aux antilles.
403 Bases climatiques de l’avertissement. Fruits 27, 665–676.

404 Ganry, J., Meyer, J., 1983. Les cerco-sporioses du bananier et leurs traitements.
405 Evolution des méthodes de traitement. Fruits 38, 3–20.

406 Glezakos, T., Moschopoulou, G., Tsigiridis, T., Kintzios, S., Yialouris, C.,
407 2010. Plant virus identification based on neural networks with evolutionary
408 preprocessing. Computers and Electronics in Agriculture 70, 263–275.

409 Hoerl, A., Kennard, R., 1988. Ridge regression. Encyclopedia of Statistical Sci-
410 ences 8, 129–136.

411 Huang, Y., Lan, Y., Thomson, S., Fang, A., Hoffmann, W., Lacey, R., 2010.
412 Development of soft computing and applications in agricultural and biological
413 engineering. Computers and Electronics in Agriculture 71 (2), 107–127.

414 Ibrahim, N., Wibowo, A., 2014. Time series support vector regression with
415 missing data treatment based variables selection for water level prediction of
416 Galas River in Kelantan Malaysia. International Journal of Applied Research
417 in Engineering and Science 3, 25–36.

418 Kim, Y., Yoo, S., Gu, Y., Lim, J., Han, D., Baik, S., 2014. Crop pests predic-
419 tion method using regression and machine learning technology: Survey. IERI
420 Procedia 6, 52–56.

421 Lukoševičius, M., 2012. A practical guide to applying echo state networks. Neu-
422 ral Networks: Tricks of the Trade, 1–20.

423 Lukoševičius, M., Jaeger, H., 2009. Reservoir computing approaches to recurrent
424 neural network training. Computer Science Review 3, 127–149.

425 Marín, D., Romero, R., Guzmán, M., Sutton, T., 2003. Black Sigatoka: An
426 increasing threat to banana cultivation. Plant Disease 87 (3), 208–222.

427 Marín Vargas, D., Romero Calderón, R., 1995. El combate de la Sigatoka Negra.
428 In: Boletín Departamento de Investigaciones. Corbana, Costa Rica, pp. 1–23.

429 McKinney, W., 2010. Data structures for statistical computing in Python. In:
430 Proceedings of the 9th Python in Science Conference. pp. 51–56.

431 Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O.,
432 Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos,
433 A., Cournapeau, D., Brucher, M., Perrot, M., Duchesnay, E., 2011. Scikit-
434 learn: Machine learning in python. Journal of Machine Learning Research 12,
435 2825–2830.

436 Press, W. H., Teukolsky, S. A., Vetterling, W. T., Flannery, B. P., 2007. Nu-
437 merical Recipes. The Art of Scientific Computing, tercera edición Edition.
438 Cambridge University Press.

439 Romero Calderón, R., 1995. Dynamics of fungicide resistant populations of *My-*
440 *cosphaerella fijiensis* and epidemiology of black Sigatoka of banana. Ph.D.
441 thesis, Department of Plant Pathology, North Carolina State University.

- 442 Rumpf, T., Mahlein, A.-K., Steiner, U., Oerke, E.-C., Dehne, H.-W., Plümer, L.,
443 2010. Early detection and classification of plant diseases with support vector
444 machines based on hyperspectral reflectance. *Computers and Electronics in*
445 *Agriculture* 74 (1), 91–99.
- 446 Smola, A. J., Schölkopf, B., 2004. A tutorial on support vector regression. *Statistics and Computing* 14, 199–222.
- 448 Soares, J., Pasqual, M., Lacerda, W., 2013. Utilization of artificial neural net-
449 works in the prediction of the bunches’ weight in banana plants. *Scientia*
450 *Horticulturae* 155, 24–29.
- 451 Soares, J., Pasqual, M., Lacerda, W., Silva, S., Donato, S., 2014. Compari-
452 son of techniques used in the prediction of yield in banana plants. *Scientia*
453 *Horticulturae* 167, 84–90.
- 454 Ternesien, E., 1985. La cercorporioses des bananiers et plantains. *Methodes de*
455 *lutte-Avertissements. Perspectives au Cameroun. Memoire de fin d’etudes.*
- 456 Tibshirani, R., 1996. Regression shrinkage and selection via the lasso. *Journal*
457 *of the Royal Statistical Society: Series B* 58 (1), 267–288.
- 458 van der Walt, S., Colbert, C., Varoquaux, G., 2011. The NumPy Array: A struc-
459 ture for efficient numerical computation. *Computing in Science & Engineering*
460 13, 22–30.
- 461 Wei, Z., Tao, T., ZhuoShu, D., Zio, E., 2013. A dynamic particle filter-support
462 vector regression method for reliability prediction. *Reliability Engineering &*
463 *System Safety*, 109–116.
- 464 Zou, H., Hastie, T., 2005. Regularization and variable selection via the elastic
465 net. *Journal of the Royal Statistical Society: Series B* 67, 301–320.